

PYTHON FOR DATA ANALYSIS PROJECT WORK

AIM : To develop a predictive model to identify the Severity of Cirrhosis Disease in an Individual based on Clinical and Laboratory Factors Using Machine Learning Model

About Columns

1. Id - Id of the patient
2. N_Days - Number of days between registration and the earlier of death, transplantation, or study analysis time in 1986
3. Status - Status of the patient: C (censored), CL (censored due to liver tx), or D (death)
4. Drug - Type of drug - D-penicillamine or placebo
5. Age - Age in days
6. Sex - M (male) or F (female)
7. Ascites - Presence of ascites: N (No) or Y (Yes)
8. Hepatomegaly - Presence of hepatomegaly: N (No) or Y (Yes)
9. Spiders - Presence of spiders: N (No) or Y (Yes)
10. Edema - Presence of edema: N (no edema and no diuretic therapy for edema), S (edema present without diuretics, or edema resolved by diuretics), or Y (edema despite diuretic therapy)
11. Bilirubin - Serum bilirubin in [mg/dl]
12. Cholesterol - Serum cholesterol in [mg/dl]
13. Albumin - Albumin in [gm/dl]
14. Copper - Urine copper in [ug/day]
15. Alk_Phos - Alkaline phosphatase in [U/liter]
16. SGOT - SGOT in [U/ml]
17. Tryglicerides - Triglycerides in [mg/dl]
18. Platelets - Platelets per cubic [ml/1000]
19. Prothrombin - Prothrombin time in seconds [s]
20. Stage - Histologic stage of disease (1, 2, or 3)

The Categorical Columns of Data are: ['Status', 'Drug', 'Sex', 'Ascites', 'Hepatomegaly', 'Spiders', 'Edema', 'Stage']

The Numerical Columns of Data are: ['N_Days', 'Age', 'Bilirubin', 'Cholesterol', 'Albumin', 'Copper', 'Alk_Phos', 'SGOT', 'Tryglicerides', 'Platelets', 'Prothrombin']

```
In [1]: #importing necessary libraries :
import pandas as pd
import numpy as np
import matplotlib
from matplotlib import pyplot as plt
import seaborn as sns
import pandas as pd
```

```
In [3]: #Loading the csv :
try:
    df = pd.read_csv("C:\\Users\\mtab\\Downloads\\cirrhosis.csv")
    print("Data set is loaded successfully...")

except FileNotFoundError as e:
    print(f"Error: {e}, Check the file path...")

except pd.errors.ParserError as e:
    print(f"Error with the file: {e}")

Data set is loaded successfully...
```

Overview of the dataset

```
In [4]: #analyzing the dataset :
df.info()
#summary of dataframes struc

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   ID                    418 non-null    int64
 1   N_Days                418 non-null    int64
 2   Status                418 non-null    object
 3   Drug                  312 non-null    object
 4   Age                   418 non-null    int64
 5   Sex                   418 non-null    object
 6   Ascites               312 non-null    object
 7   Hepatomegaly          312 non-null    object
 8   Spiders               312 non-null    object
 9   Edema                 418 non-null    object
10  Bilirubin              418 non-null    float64
11  Cholesterol            284 non-null    float64
12  Albumin                418 non-null    float64
13  Copper                 310 non-null    float64
14  Alk_Phos              312 non-null    float64
15  SGOT                   312 non-null    float64
16  Tryglicerides          282 non-null    float64
17  Platelets              407 non-null    float64
18  Prothrombin            416 non-null    float64
19  Stage                  412 non-null    float64
dtypes: float64(10), int64(3), object(7)
memory usage: 65.4+ KB
```

```
In [5]: df.shape
```

```
Out[5]: (418, 20)
```

```
In [6]: #dataset has 418 records and 20 fields
```

In [7]: *#reading the first and last 5 records of the cirrhosis dataset :*
`df.head(5)`

Out[7]:

	ID	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin
0	1	400	D	D- penicillamine	21464	F	Y	Y	Y	Y	14.5	261.0	2.60
1	2	4500	C	D- penicillamine	20617	F	N	Y	Y	N	1.1	302.0	4.14
2	3	1012	D	D- penicillamine	25594	M	N	N	N	S	1.4	176.0	3.48
3	4	1925	D	D- penicillamine	19994	F	N	Y	Y	S	1.8	244.0	2.54
4	5	1504	CL	Placebo	13918	F	N	Y	Y	N	3.4	279.0	3.53

In [8]: `df.tail(5)`

Out[8]:

	ID	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Co
413	414	681	D	NaN	24472	F	NaN	NaN	NaN	N	1.2	NaN	2.96	
414	415	1103	C	NaN	14245	F	NaN	NaN	NaN	N	0.9	NaN	3.83	
415	416	1055	C	NaN	20819	F	NaN	NaN	NaN	N	1.6	NaN	3.42	
416	417	691	C	NaN	21185	F	NaN	NaN	NaN	N	0.8	NaN	3.75	
417	418	976	C	NaN	19358	F	NaN	NaN	NaN	N	0.7	NaN	3.29	

In [9]: *#checking for any null values and handling them :*

`df.isnull().sum()` *#null values observed in 8 columns*

Out[9]:

```
ID          0
N_Days      0
Status      0
Drug        106
Age         0
Sex         0
Ascites     106
Hepatomegaly 106
Spiders     106
Edema       0
Bilirubin   0
Cholesterol 134
Albumin     0
Copper      108
Alk_Phos    106
SGOT        106
Tryglicerides 136
Platelets   11
Prothrombin 2
Stage       6
dtype: int64
```

```
In [10]: #filling null values using median and mode
# mode --> categorical values
# median --> numerical values

df['Ascites'] = df['Ascites'].fillna(df['Ascites'].mode()[0])
df['Hepatomegaly'] = df['Hepatomegaly'].fillna(df['Hepatomegaly'].mode()[0])
df['Spiders'] = df['Spiders'].fillna(df['Spiders'].mode()[0])
df['Drug'] = df['Drug'].fillna(df['Drug'].mode()[0])
df['Stage'] = df['Stage'].fillna(df['Stage'].mode()[0])

df['Cholesterol'] = df['Cholesterol'].fillna(df['Cholesterol'].median())
df['Copper'] = df['Copper'].fillna(df['Copper'].median())
df['SGOT'] = df['SGOT'].fillna(df['SGOT'].median())
df['Tryglicerides'] = df['Tryglicerides'].fillna(df['Tryglicerides'].median())
df['Platelets'] = df['Platelets'].fillna(df['Platelets'].median())
df['Prothrombin'] = df['Prothrombin'].fillna(df['Prothrombin'].median())
df['Alk Phos'] = df['Alk Phos'].fillna(df['Alk Phos'].median())
```

```
In [11]: #rechecking for null values
df.isnull().sum()
```

```
Out[11]: ID                0
N_Days                0
Status                0
Drug                 0
Age                  0
Sex                  0
Ascites              0
Hepatomegaly         0
Spiders              0
Edema                0
Bilirubin            0
Cholesterol           0
Albumin              0
Copper               0
Alk_Phos             0
SGOT                 0
Tryglicerides        0
Platelets            0
Prothrombin          0
Stage                0
dtype: int64
```

```
In [12]: #Looking for any duplicate rows and dropping if any
num_duplicates = df.duplicated().sum()
print(f'Number of duplicate rows: {num_duplicates}')


Number of duplicate rows: 0
```

```
In [13]: df['Age'] = df['Age']/365 #converting age(given-in days) into years
```

Visual Analysis of the Dataset

Distribution of Gender

```
In [14]: plt.figure(figsize=(6, 6))  
sns.countplot(x='Sex', data=df)  
plt.title('Distribution of Gender', fontsize=16)  
plt.xlabel('Gender')  
plt.ylabel('Count')  
plt.show()
```




Interpretation : Among the total number of individuals diagnosed with cirrhosis, approximately 375-380 are female patients and 40-43 male patients. This indicates a greater occurrence of cirrhosis in females as compared to males.

Distribution By Age

```
In [15]: plt.hist(df['Age'], bins=20, color='purple', edgecolor='pink')

plt.title('Distribution by Age', fontsize=12)
plt.xlabel('Age', fontsize=12)
plt.ylabel('Count', fontsize=12)
plt.show()
```

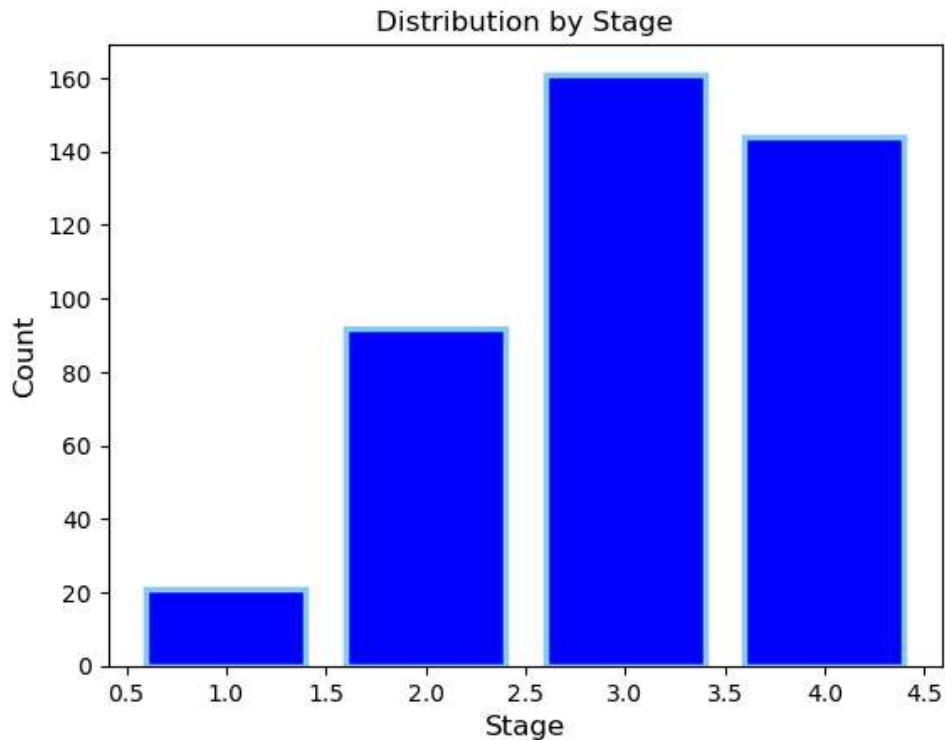


Interpretation : From the above graph we can say that there is a gradual increase in the number of people suffering from cirrhosis as the age of an individual increases from 30-50 years of age, following which the number of people with cirrhosis starts to go down.

count of people vs the stage of cirrhosis

```
In [16]: stage_counts = df['Stage'].value_counts()


plt.bar(stage_counts.index, stage_counts.values, color='blue', edgecolor='skyblue',linewidth=2.5)
plt.title('Distribution by Stage', fontsize=12)
plt.xlabel('Stage', fontsize=12)
plt.ylabel('Count', fontsize=12)
plt.show()
```



Interpretation : Among the total number of people, the highest number of cases is observed in stage 3 cirrhosis with a count of approximately 160 individuals, followed by around 150 people suffering from stage 4 cirrhosis. The count for stage 2 cirrhosis is relatively low, with a record of around 90 people, while only 20 people are recorded in stage 1 cirrhosis.

Relationship between Age vs Stage

```
In [17]: #Age vs Stage
plt.figure(figsize=(8, 5))
sns.boxplot(data=df, x='Stage', y='Age', palette='Set2')
plt.title('Age vs Stage')
plt.xlabel('Stage')
plt.ylabel('Age')
plt.show()
```



Stage 1: The ages in Stage 1 are mostly concentrated between 30 and 50, with a few extending to around 60.

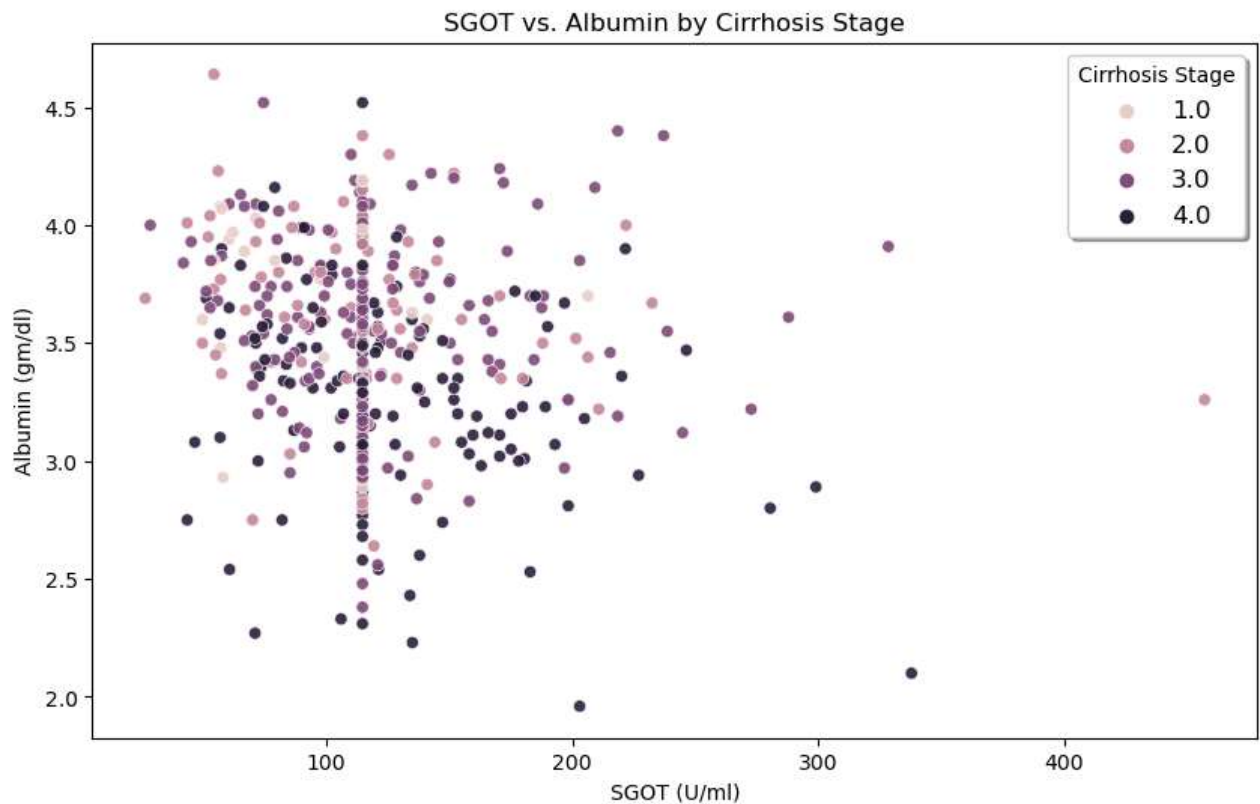
Stage 2: The range of ages in Stage 2 is similar to Stage 1, but the median age seems slightly higher (around 48).

Stage 3: In Stage 3, the median age is slightly higher than in the previous stages. The range of ages is also wider, with a few individuals between the age of 60-70 years.

Stage 4: Stage 4 shows the highest median age among all stages (around 53-54 years). The range is also quite wide, suggesting that individuals at this stage can be quite diverse in terms of age.

SGOT VS ALBUMIN by Cirrhosis Stage

```
In [18]: #SGOT vs Albumin
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='SGOT', y='Albumin', hue='Stage', alpha=0.9)
plt.title('SGOT vs. Albumin by Cirrhosis Stage')
plt.xlabel('SGOT (U/ml)')
plt.ylabel('Albumin (gm/dl)')
plt.legend(title='Cirrhosis Stage', loc="best", shadow=True, fontsize="large")
plt.show()
```



There is a negative correlation between SGOT and Albumin levels, As SGOT levels increase, indicating greater liver damage, albumin levels tend to decrease.

The scatter plot shows that individuals with more advanced cirrhosis stages (3.0 and 4.0) generally have lower albumin levels compared to those in earlier stages (1.0 and 2.0).

pie chart plot

```
In [19]: import pandas as pd
import matplotlib.pyplot as plt

fig, axs = plt.subplots(2, 2, figsize=(16, 8))

# Pie chart for Hepatomegaly
hepatomegaly_counts = df['Hepatomegaly'].value_counts()
axs[0, 0].pie(hepatomegaly_counts, labels=hepatomegaly_counts.index,
              autopct='%1.1f%%', colors=['purple', 'cyan'], explode=[0, 0.05])
axs[0, 0].set_title('Proportion of Patients with Hepatomegaly')

# Pie chart for Edema
edema_counts = df['Edema'].value_counts()
axs[0, 1].pie(edema_counts, labels=edema_counts.index,
              autopct='%1.1f%%', colors=['pink', 'purple', 'cyan'], explode=[0, 0.1, 0.2])
axs[0, 1].set_title('Proportion of Patients with Edema')


# Pie chart for Drug
drug_counts = df['Drug'].value_counts()
axs[1, 0].pie(drug_counts, labels=drug_counts.index,
              autopct='%1.1f%%', colors=['pink', 'cyan'], explode=[0, 0.05])
axs[1, 0].set_title('Drug Distribution')

# Pie chart for Ascites
ascites_counts = df['Ascites'].value_counts()
axs[1, 1].pie(ascites_counts, labels=ascites_counts.index,
              autopct='%1.1f%%', colors=['purple', 'cyan'], explode=[0, 0.05])
axs[1, 1].set_title('Ascites Distribution')


# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()

#Edema - Presence of edema: N (no edema and no diuretic therapy for edema),
#S (edema present without diuretics, or edema resolved by diuretics), or Y (edema despite diuretic
```

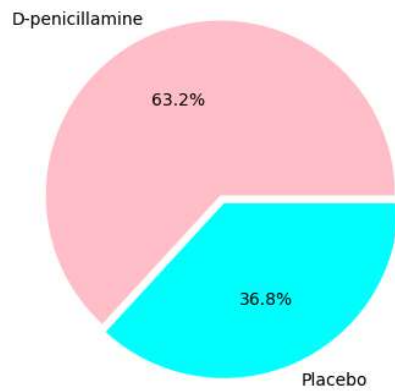
Proportion of Patients with Hepatomegaly




Proportion of Patients with Edema



Drug Distribution



Ascites Distribution



Plot for Average Bilirubin Levels by Cirrhosis Stage

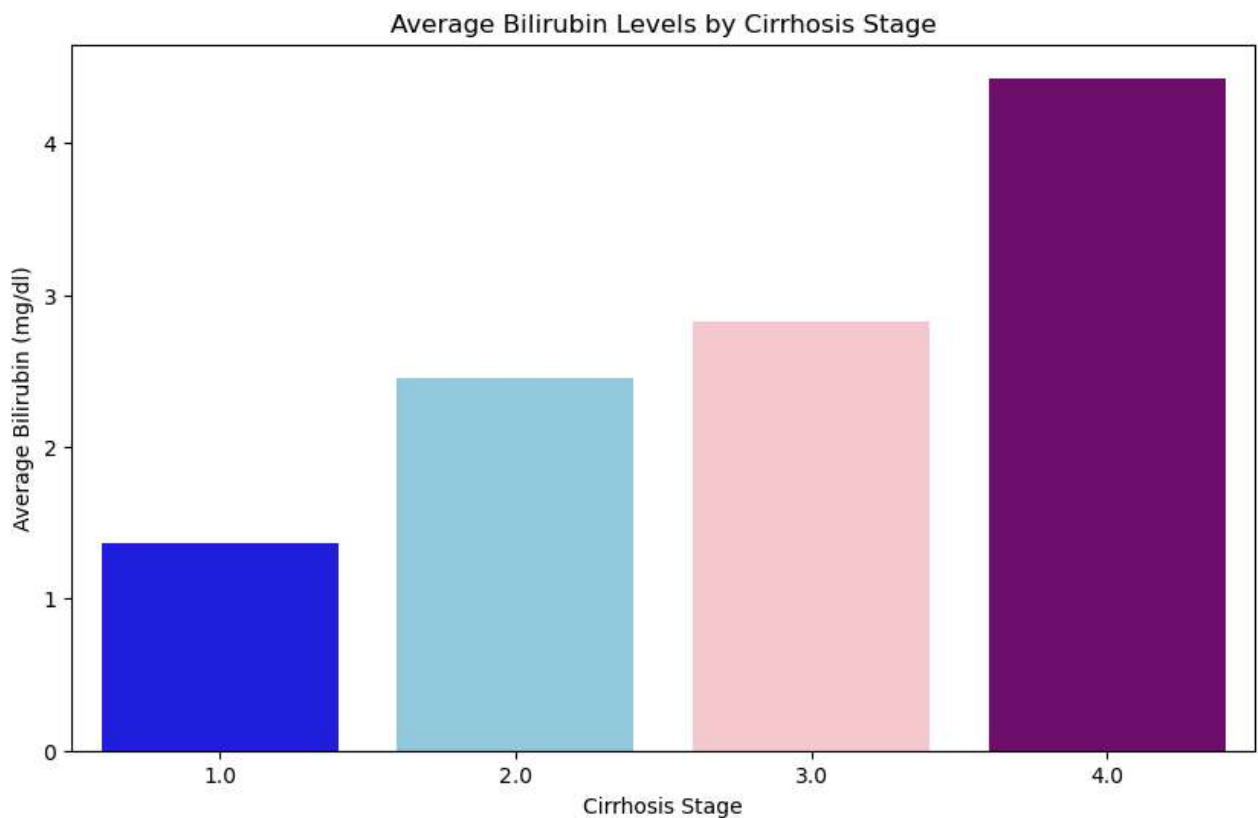
```
In [20]: clrs = ['blue', 'skyblue', 'pink', 'purple']

plt.figure(figsize=(10, 6))
sns.barplot(data=df, x='Stage', y='Bilirubin', ci=None, palette=clrs)
plt.title('Average Bilirubin Levels by Cirrhosis Stage')
plt.xlabel('Cirrhosis Stage')
plt.ylabel('Average Bilirubin (mg/dl)')
plt.show()
```

C:\Users\mtab\AppData\Local\Temp\ipykernel_11292\3671046248.py:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=df, x='Stage', y='Bilirubin', ci=None, palette=clrs)
```




Bilirubin levels are relatively low in this early stage of cirrhosis. While the levels reach their highest point in Stage 4, indicating significant liver damage.

Pair plot

```
In [21]: # Selecting relevant numerical columns along with the Stage column
s_Cols = ['Cholesterol', 'Albumin', 'Platelets', 'Stage']
subset_data = df[s_Cols]

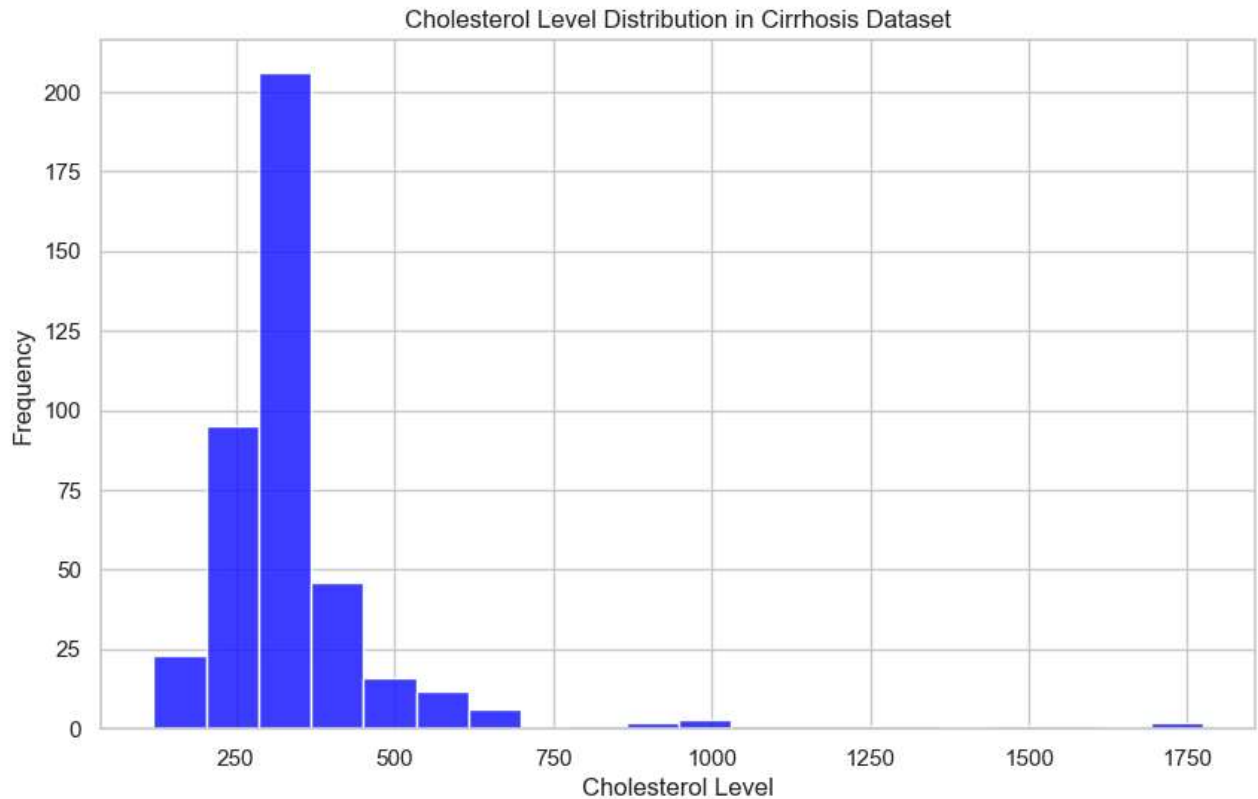
sns.pairplot(subset_data, hue='Stage')
plt.show()
```



```
In [22]: #This graph shows how Cholesterol, Albumin, and Platelet Levels change
#as the disease (cirrhosis) progresses. We can see that there are some weak relationships between the
#but they don't tell us much about the severity of the disease.
```

```
In [23]: sns.set(style="whitegrid")

# Plotting the distribution of cholesterol levels
plt.figure(figsize=(10, 6))
sns.histplot(df['Cholesterol'], bins=20, color='blue')
plt.title('Cholesterol Level Distribution in Cirrhosis Dataset')
plt.xlabel('Cholesterol Level')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()
```



```
In [24]: # Defining the cholesterol range # as from the above plot we can say max. people suffering
# from cirrhosis have cholesterol level btw 250-500
lower_th = 250
upper_th = 500

# Creating a boolean mask for cholesterol levels
m = (df['Cholesterol'] >= lower_th) & (df['Cholesterol'] <= upper_th)

# Filter the DataFrame based on the mask
filtered_df = df[m]
num_rows = filtered_df.shape[0]
print(f"\nNumber of patients with Cholesterol between {lower_th} and {upper_th}: {num_rows}")
```

Number of patients with Cholesterol between 250 and 500: 311

```
In [25]: #viewing the first 5 row of the above
print("\n\nFiltered Rows (Cholesterol between 250 and 500):")
print(filtered_df.iloc[:5])
```

Filtered Rows (Cholesterol between 250 and 500):

	ID	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	\
0	1	400	D	D-penicillamine	58.805479	F	Y		Y
1	2	4500	C	D-penicillamine	56.484932	F	N		Y
4	5	1504	CL	Placebo	38.131507	F	N		Y
6	7	1832	C	Placebo	55.572603	F	N		Y
7	8	2466	D	Placebo	53.093151	F	N		N

	Spiders	Edema	Bilirubin	Cholesterol	Albumin	Copper	Alk_Phos	SGOT	\
0	Y	Y	14.5	261.0	2.60	156.0	1718.0	137.95	
1	Y	N	1.1	302.0	4.14	54.0	7394.8	113.52	
4	Y	N	3.4	279.0	3.53	143.0	671.0	113.15	
6	N	N	1.0	322.0	4.09	52.0	824.0	60.45	
7	N	N	0.3	280.0	4.00	52.0	4651.2	28.38	

	Tryglicerides	Platelets	Prothrombin	Stage
0	172.0	190.0	12.2	4.0
1	88.0	221.0	10.6	3.0
4	72.0	136.0	10.9	3.0
6	213.0	204.0	9.7	3.0
7	189.0	373.0	11.0	3.0

```
In [26]: # Counting occurrences of each stage with filtered cholesterol levels :
stage_counts = filtered_df['Stage'].value_counts()

print("\nStage Counts:")
print(stage_counts)
```

Stage Counts:

3.0 121

4.0 106

2.0 73

1.0 11

Name: Stage, dtype: int64

Pushing data into database

```
In [ ]: import sqlite3

conn = sqlite3.connect('Project_PYDA.db')

df.to_sql('CirrhosisData', conn, if_exists='replace', index=False)

conn.close()

print("DataFrame pushed to SQLite database successfully!")
DataFrame pushed to SQLite database successfully!
```

SQLiteStudio (3.4.4) - [CirrhosisData (Assig_project_PYDA)]

DatabaseStructureViewToolsHelp

Databases

Assig_project_PYDA (SQLite 3)

Tables (1)

CirrhosisData

Views

StructureDataConstraintsIndexesTriggersDDL

Grid viewForm view

Filter dataTotal rows loaded: 418

ID	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomeg	Spiders	Edema	Bilirubin	Cholestero	Albumin	Copper	Alk_Phos	
403	403	1141	C	D-penicillamine	66.04657534246576	F	N	Y	N	N	2.5	309.5	3.33	73	1259
404	404	1092	C	D-penicillamine	40.02739726027397	F	N	Y	N	N	4.6	309.5	3.6	73	1259
405	405	1150	C	D-penicillamine	52.03561643835616	F	N	Y	N	N	1	309.5	3.64	73	1259
406	406	703	D	D-penicillamine	46.03287671232876	F	N	Y	N	N	4.5	309.5	2.68	73	1259
407	407	1129	C	D-penicillamine	54.03835616438356	M	N	Y	N	N	1.1	309.5	3.69	73	1259
408	408	1086	C	D-penicillamine	51.03561643835616	F	N	Y	N	S	1.9	309.5	3.17	73	1259
409	409	1067	C	D-penicillamine	43.03013698630137	F	N	Y	N	N	0.7	309.5	3.73	73	1259
410	410	1072	C	D-penicillamine	39.02739726027397	F	N	Y	N	N	1.5	309.5	3.81	73	1259
411	411	1119	C	D-penicillamine	51.03561643835616	F	N	Y	N	N	0.6	309.5	3.57	73	1259
412	412	1097	C	D-penicillamine	67.04657534246576	F	N	Y	N	N	1	309.5	3.58	73	1259
413	413	989	C	D-penicillamine	35.02465753424657	F	N	Y	N	N	0.7	309.5	3.23	73	1259
414	414	681	D	D-penicillamine	67.04657534246576	F	N	Y	N	N	1.2	309.5	2.96	73	1259
415	415	1103	C	D-penicillamine	39.02739726027397	F	N	Y	N	N	0.9	309.5	3.83	73	1259
416	416	1055	C	D-penicillamine	57.03835616438356	F	N	Y	N	N	1.6	309.5	3.42	73	1259
417	417	691	C	D-penicillamine	58.04109589041096	F	N	Y	N	N	0.8	309.5	3.75	73	1259
418	418	976	C	D-penicillamine	53.03561643835616	F	N	Y	N	N	0.7	309.5	3.29	73	1259

SQLiteStudio (3.4.4) - [CirrhosisData (Assig_project_PYDA)]

Database Structure View Tools Help

Databases

Filter by name

Assig_project_PYDA (SQLite 3)

Tables (1)

CirrhosisData

Views

Grid view Form view

Filter data Total rows loaded: 418

ID	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomeg	Spiders	Edema	Bilirubin	Cholestero	Albumin	Copper	Alk_Phos	
1	1	400	D	D-penicillamine	58.8054794520548	F	Y	Y	Y	Y	14.5	261	2.6	156	1718
2	2	4500	C	D-penicillamine	56.48493150684931	F	N	Y	Y	N	1.1	302	4.14	54	7394.8
3	3	1012	D	D-penicillamine	70.12054794520547	M	N	N	N	S	1.4	176	3.48	210	516
4	4	1925	D	D-penicillamine	54.77808219178083	F	N	Y	Y	S	1.8	244	2.54	64	6121.8
5	5	1504	CL	Placebo	38.13150684931507	F	N	Y	Y	N	3.4	279	3.53	143	671
6	6	2503	D	Placebo	66.30410958904109	F	N	Y	N	N	0.8	248	3.98	50	944
7	7	1832	C	Placebo	55.57260273972603	F	N	Y	N	N	1	322	4.09	52	824
8	8	2466	D	Placebo	53.09315068493151	F	N	N	N	N	0.3	280	4	52	4651.2
9	9	2400	D	D-penicillamine	42.53698630136986	F	N	N	Y	N	3.2	562	3.08	79	2276
10	10	51	D	Placebo	70.6082191780822	F	Y	N	Y	Y	12.6	200	2.74	140	918
11	11	3762	D	Placebo	53.75068493150685	F	N	Y	Y	N	1.4	259	4.16	46	1104
12	12	304	D	Placebo	59.17808219178082	F	N	N	Y	N	3.6	236	3.52	94	591
13	13	3577	C	Placebo	45.72054794520548	F	N	N	N	N	0.7	281	3.85	40	1181
14	14	1217	D	Placebo	56.26027397260274	M	Y	Y	N	Y	0.8	309.5	2.27	43	728
15	15	3584	D	D-penicillamine	64.69041095890411	F	N	N	N	N	0.8	231	3.87	173	9009.8
16	16	3672	C	Placebo	40.47123287671233	F	N	N	N	N	0.7	204	3.66	28	685

Model Building

```
In [27]: # Label encoding for categorical features
df['Status'] = df['Status'].map({'C': 0, 'CL': 1, 'D': 2})
df['Drug'] = df['Drug'].map({'D-penicillamine': 0, 'Placebo': 1})
df['Sex'] = df['Sex'].map({'F': 0, 'M': 1})
df['Ascites'] = df['Ascites'].map({'N': 0, 'Y': 1})
df['Hepatomegaly'] = df['Hepatomegaly'].map({'N': 0, 'Y': 1})
df['Spiders'] = df['Spiders'].map({'N': 0, 'Y': 1})
df['Edema'] = df['Edema'].map({'N': 0, 'S': 1, 'Y': 2})
```



```
In [32]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

X = df.drop(['ID', 'Status', 'Drug', 'N_Days'], axis=1) # Using all features except these
y = df['Status']

# Scaling the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
# fits the scaler to X and transforms it, resulting in X_scaled, where each feature will
# have a mean of 0 and a standard deviation of 1.python

# Splitting the data into training and testing sets - (80:20)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

Logistic Regression :

```
In [33]: model = LogisticRegression(max_iter=5000, random_state=42) #A logistic regression model is created

# Training the model
model.fit(X_train, y_train)

# Making predictions on the test set
y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy (train-test split): {accuracy:.2f}')

print("Classification Report (train-test split):")
print(classification_report(y_test, y_pred))

cm1 = confusion_matrix(y_test, y_pred)
print("Confusion_Matrix : \n",cm1)
```

```
Accuracy (train-test split): 0.80
Classification Report (train-test split):
```

	precision	recall	f1-score	support
0	0.78	0.86	0.82	44
1	0.00	0.00	0.00	4
2	0.83	0.81	0.82	36
accuracy			0.80	84
macro avg	0.53	0.56	0.54	84
weighted avg	0.76	0.80	0.78	84

```
Confusion_Matrix :
[[38  0  6]
 [ 4  0  0]
 [ 7  0 29]]
```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```


C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [34]: # Creating a heatmap for the confusion matrix for models
plt.figure(figsize=(8, 6))
sns.heatmap(cm1, annot=True, fmt='d', cmap='Purples',
            xticklabels=['Predicted Class 0', 'Predicted Class 1'],
            yticklabels=['Actual Class 0', 'Actual Class 1'])
plt.title('Confusion Matrix Heatmap')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



Random Forest :

```
In [35]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

model2 = RandomForestClassifier()
model2.fit(X_train, y_train)

#accuracy calculation :
y_train_pred = model2.predict(X_train)
training_accuracy = accuracy_score(y_train, y_train_pred)

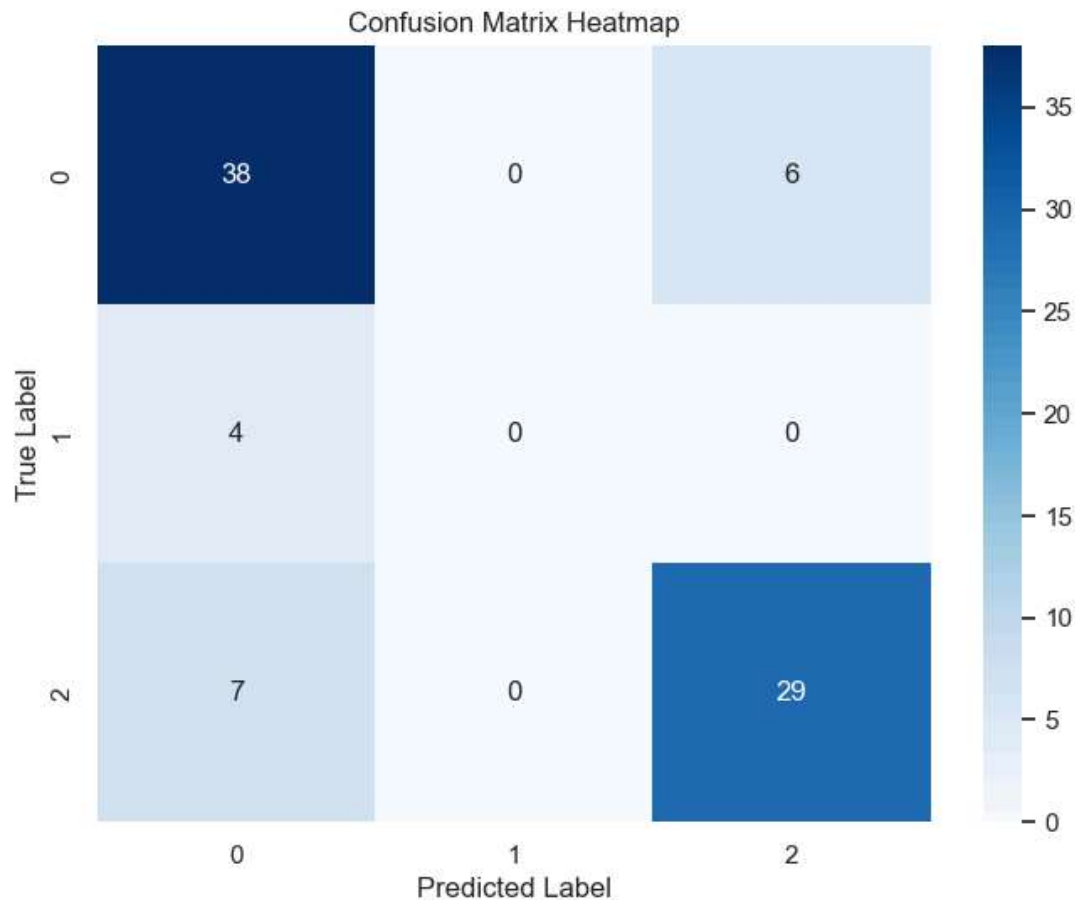
print("Training Accuracy:", training_accuracy)

#Calculating test accuracy
y_test_pred = model2.predict(X_test)
test_accuracy = accuracy_score(y_test, y_test_pred)
print("Test Accuracy:", test_accuracy)

cm_test = confusion_matrix(y_test, y_test_pred)
print("Confusion Matrix (test set):\n", cm_test)

Training Accuracy: 1.0
Test Accuracy: 0.7976190476190477
Confusion Matrix (test set):
[[38  0  6]
 [ 4  0  0]
 [ 7  0 29]]
```

```
In [36]: plt.figure(figsize=(8, 6))
sns.heatmap(cm_test, annot=True, fmt='d', cmap='Blues', xticklabels=model2.classes_, yticklabels=model2.classes_)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix Heatmap')
plt.show()
```



Conclusion :

After evaluating the performance of two machine learning models - Logistic Regression, and Random Forest on the Cirrhosis dataset, Logistic Regression emerged as the most suitable model with an accuracy of about 80%.

In []: